

# Memory Smearing: Myth or Reality?

---

Fabio Pagani



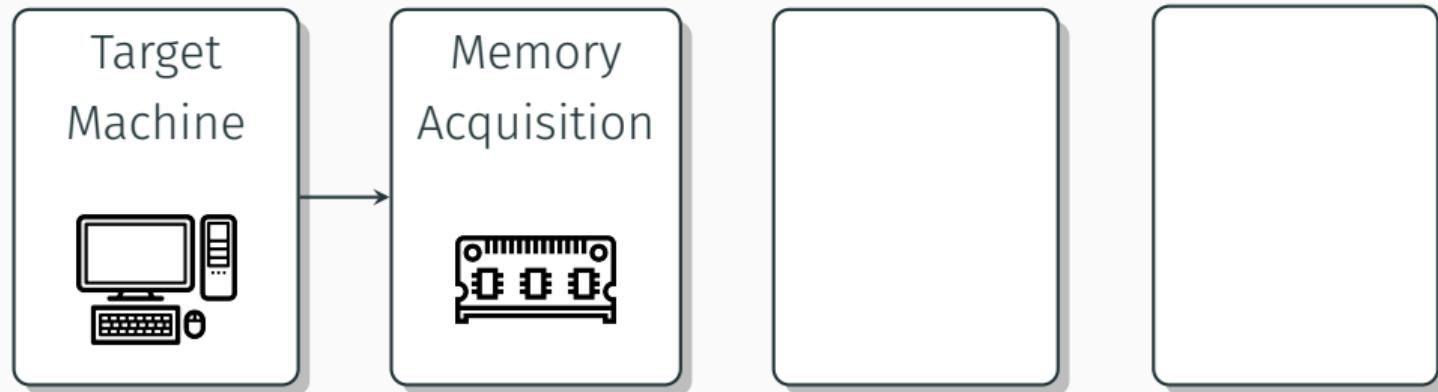
30th September 2019

# Memory Forensics - Introduction

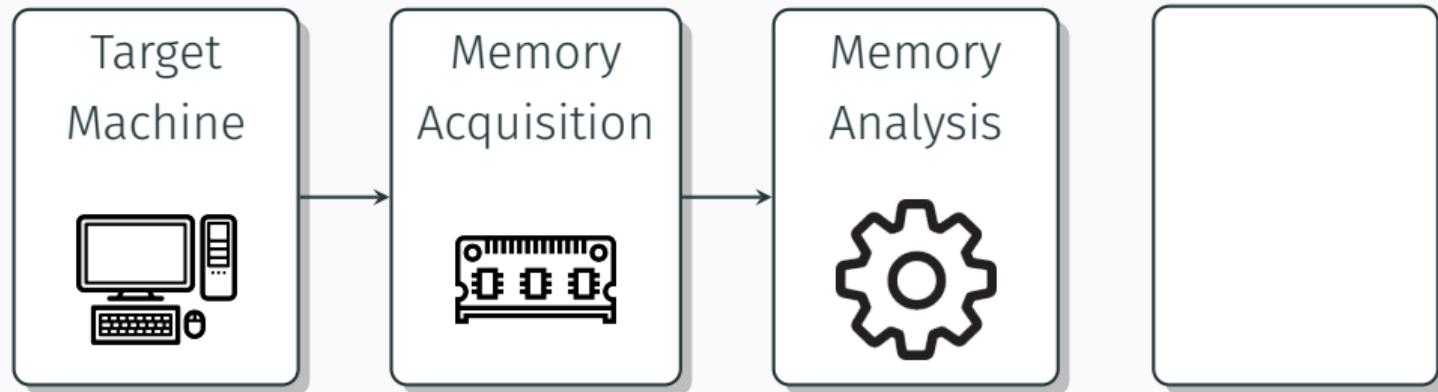
Target  
Machine



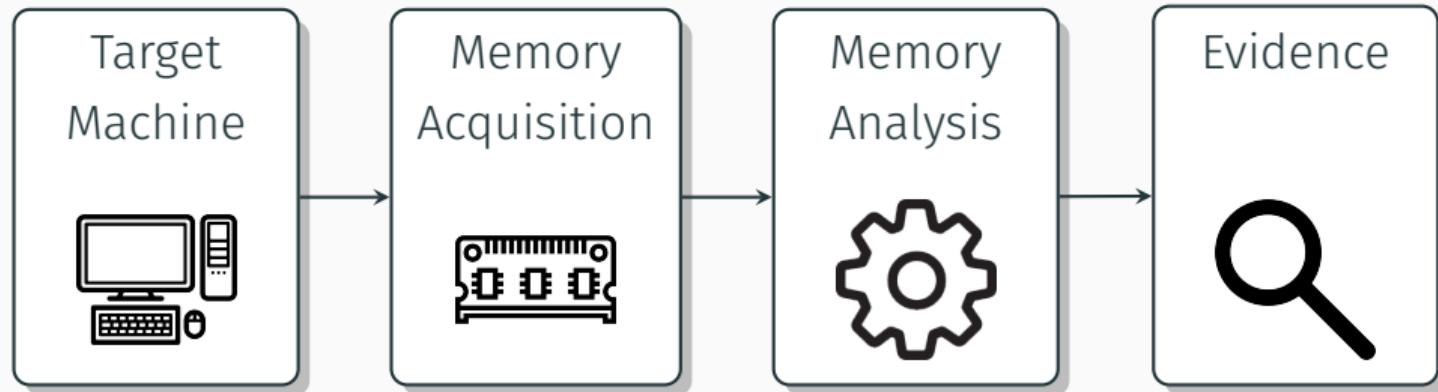
# Memory Forensics - Introduction



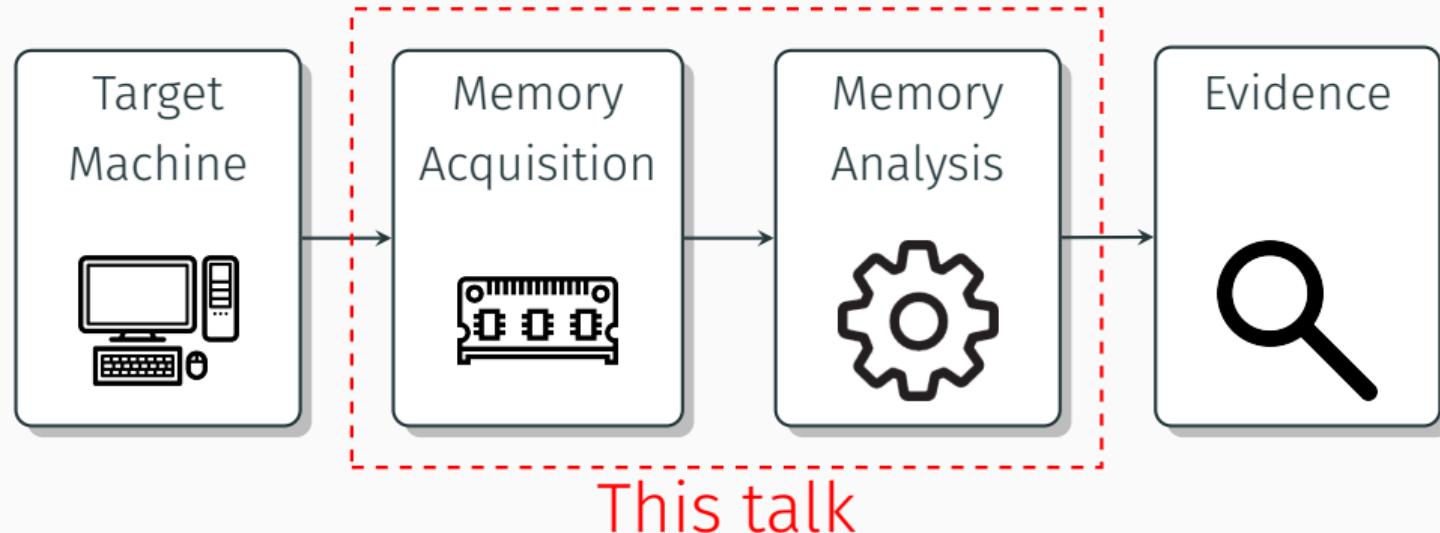
# Memory Forensics - Introduction



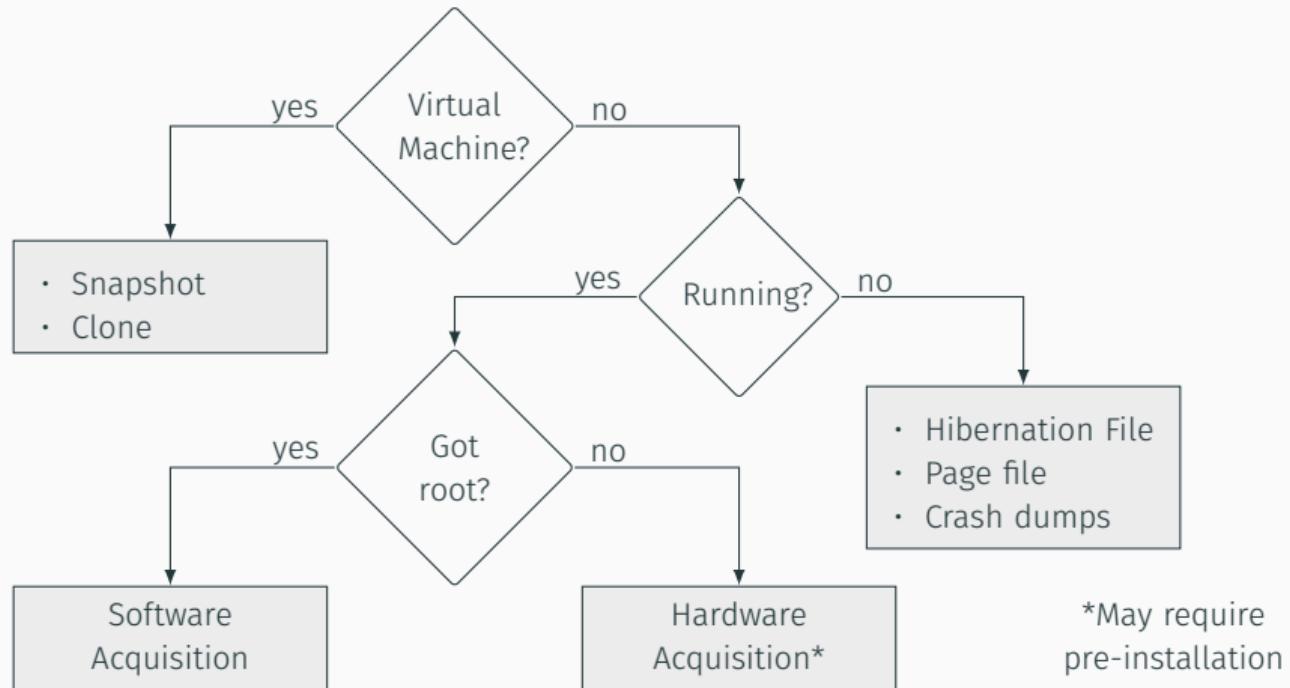
# Memory Forensics - Introduction



# Memory Forensics - Introduction

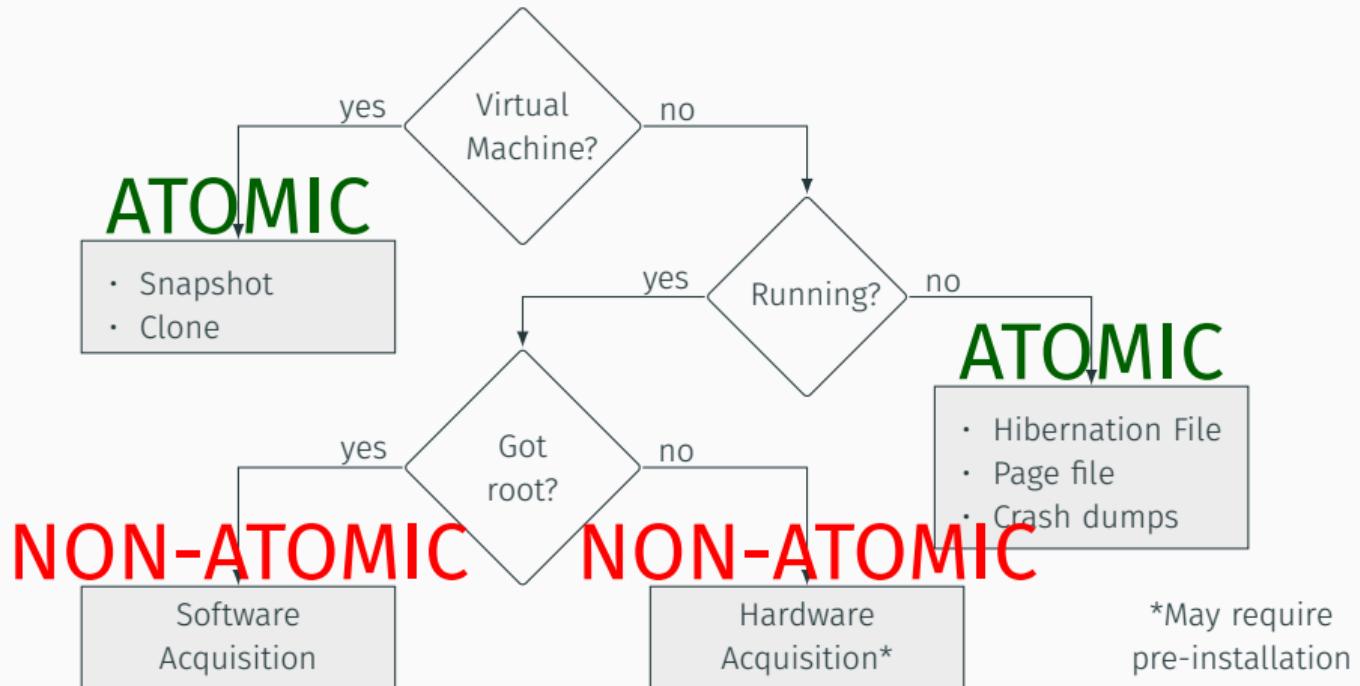


# Memory Acquisition



*Decision tree adapted from The Art of Memory Forensics*

# Memory Acquisition



*Decision tree adapted from The Art of Memory Forensics*

# ATOMIC



**ATOMIC**



**NON-ATOMIC**



# Memory Analysis

- The “core” of memory forensics.
- Several frameworks: Volatility, Rekall (Google), Mandiant’s Memoryze..
- Examples of information that can be extracted:
  - Processes → list/tree, open files, memory mappings, extract executable and shared libraries
  - Modules → list, code, unloaded modules
  - Networking → connections, sockets, arp table
  - Windows Registry → keys, password hashes
  - System information → clipboard content, screenshot
- Every task is “organized” in a plugin

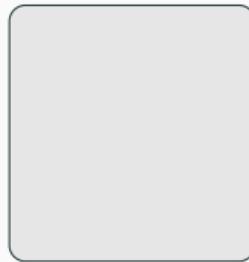
# Memory Analysis

task\_struct

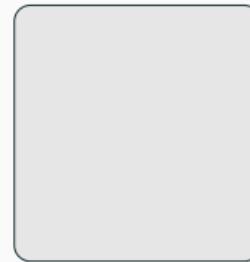


init\_task

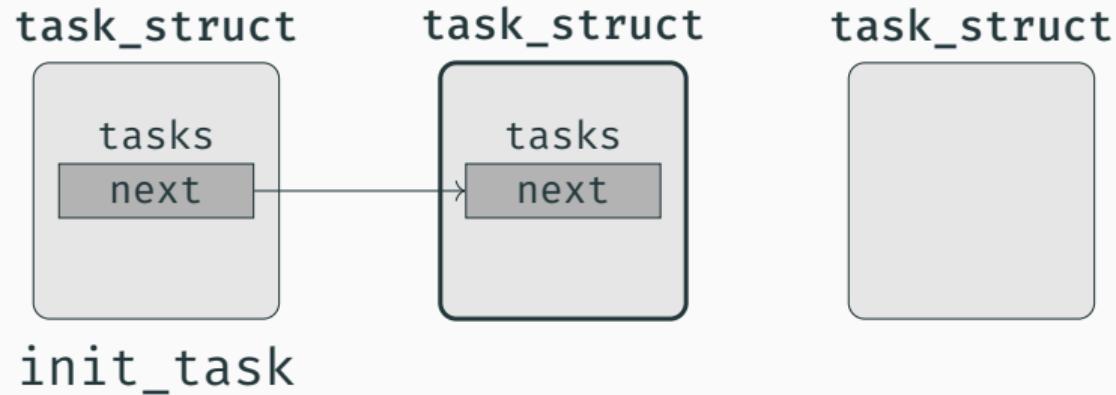
task\_struct



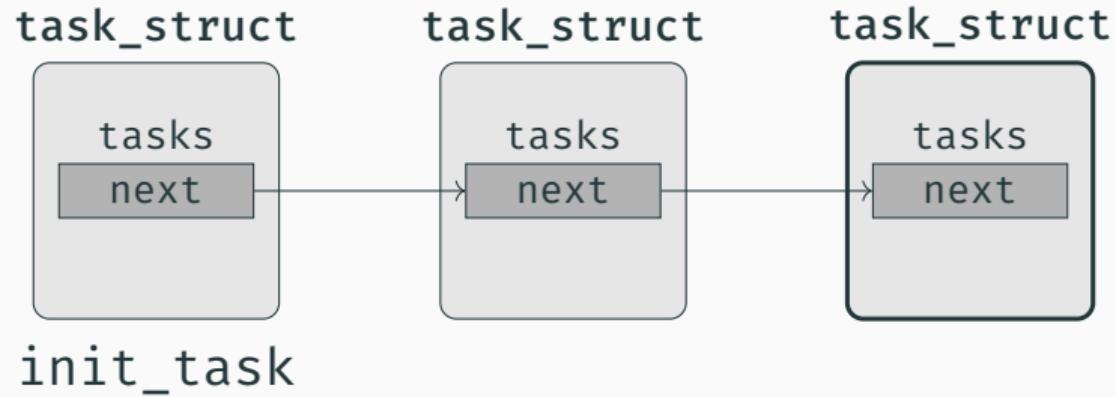
task\_struct



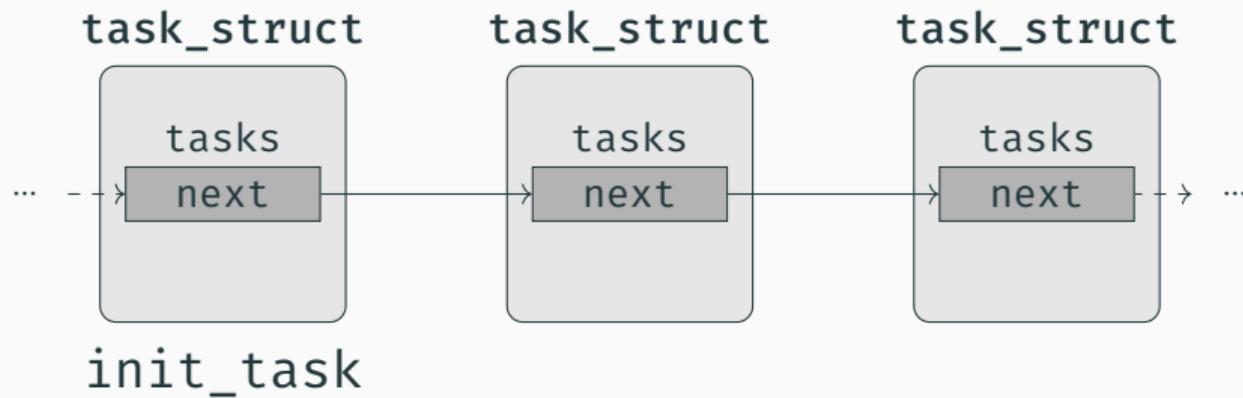
# Memory Analysis



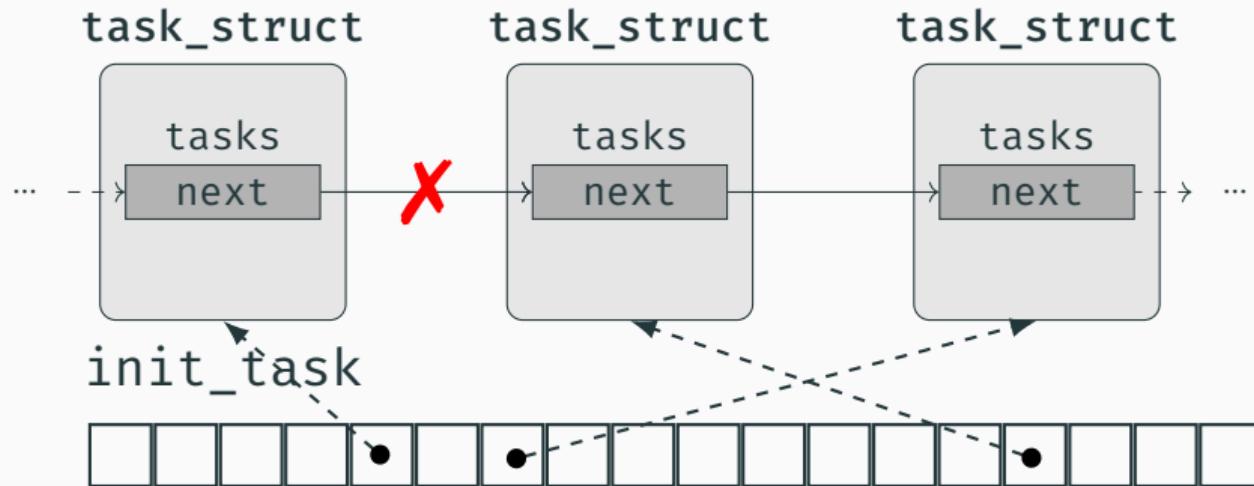
# Memory Analysis



# Memory Analysis



# Memory Analysis



## Problem

Some pointers can be **inconsistent!**

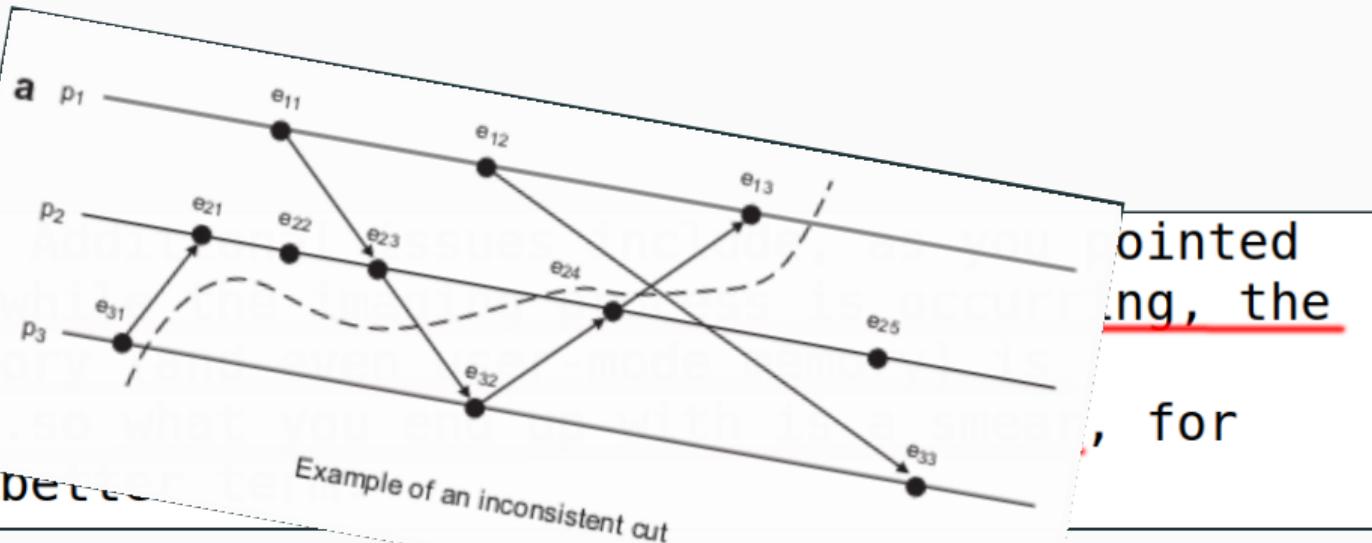
## Memory Smearing - History

pagefile. Additional issues include, as you pointed out, that while the imaging process is occurring, the kernel memory (and even user-mode memory) is changing...so what you end up with is a smear, for want of a better term.

Alan Carvey – Security Incidents ML (2005)

## Memory Smearing - History

pagefile.out, that kernel memory changing want of a better



ointed  
ng, the  
, for

Vomel et. al – Correctness, atomicity, and integrity: Defining criteria for forensically-sound memory acquisition (DFRWS 2009)

## Memory Smearing - History

pagefi  
out, t  
kernel  
changi  
want o

In about every fifth memory dump acquired via kernel-level acquisition we were confronted with inconsistent page tables. While almost the whole virtual address space of our payload application RAMMANGLE.EXE could be reconstructed, a few pages were sporadically mismapped to virtual memory of other processes, unused physical memory or kernel memory. The reason for this is yet unknown to us, however, because all tested kernel-level acquisition tools exhibited the same behavior, regardless of the acquisition method (either using `MmMapIoSpace()`, the `\Device\PhysicalMemory` device or PTE remapping) we do not consider it to be a tool error. However, on the

nted  
, the  
or

Gruhn et. al — Evaluating atomicity, and integrity of correct memory acquisition methods (DFRWS 2016)

## Memory Smearing - History

In about every fifth memory dump acquired via low-level acquisition we were confronted with page tables. While this was a problem of page tables, it was also a problem of page smearing.

page  
out,  
kern  
chan  
want

Current issues – page smearing

The following sections describe current approaches to acquisition across all major operating systems, along with the limitations of these approaches. Each section is structured to describe the state-of-the-art, its limitations, and future directions to improve acquisition techniques and procedures. We start with page smearing as it is one of the most pressing issues.

ted  
the

Case and Richard – Memory forensics: The path forward (DFWRS 2017)

## Memory Smearing - History

- To have a good memdump we need to freeze the OS

*But if you do you can have some troubles*

- You can trigger a blue screen (not really cool)

In some cases (10% to 20%), volatility and windbg can't analyze them. You can't always just make another dump.

page  
out,  
kern  
change  
want

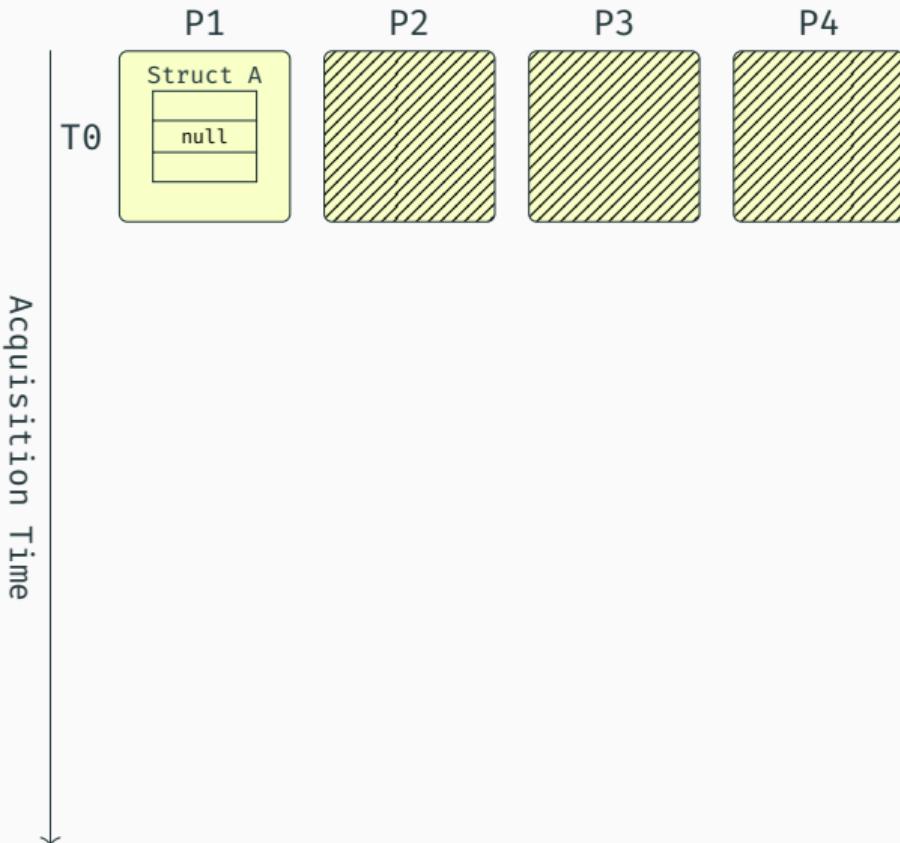
Current  
The  
tion  
of th  
of t  
situ  
is

ted  
the

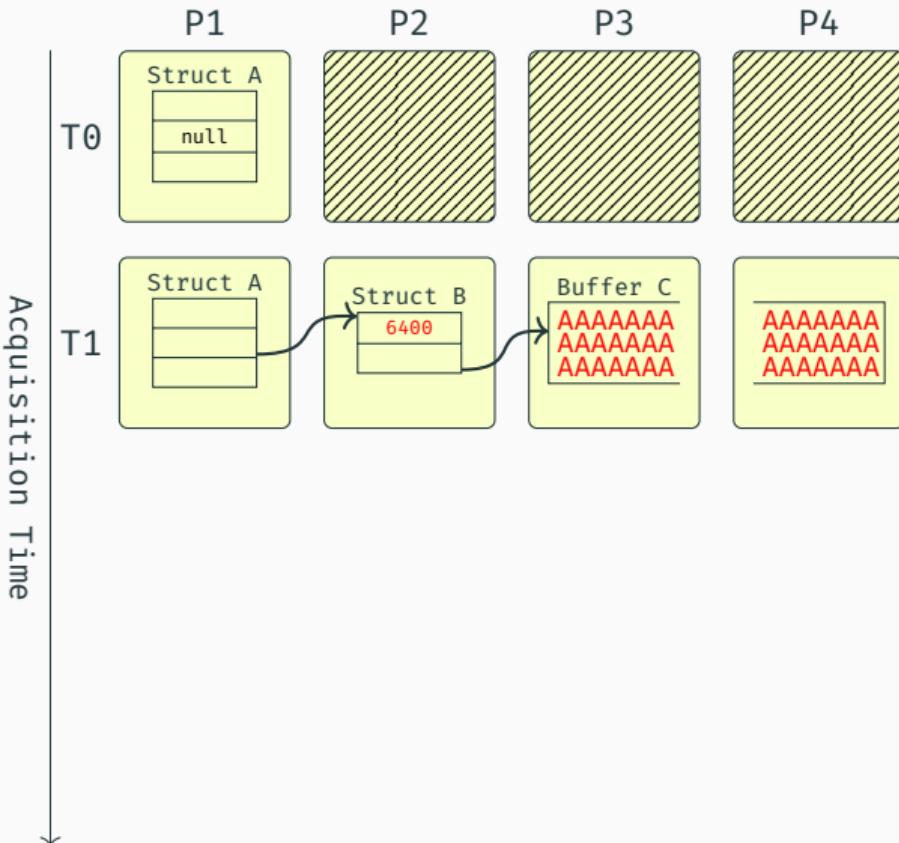
## The Problem

---

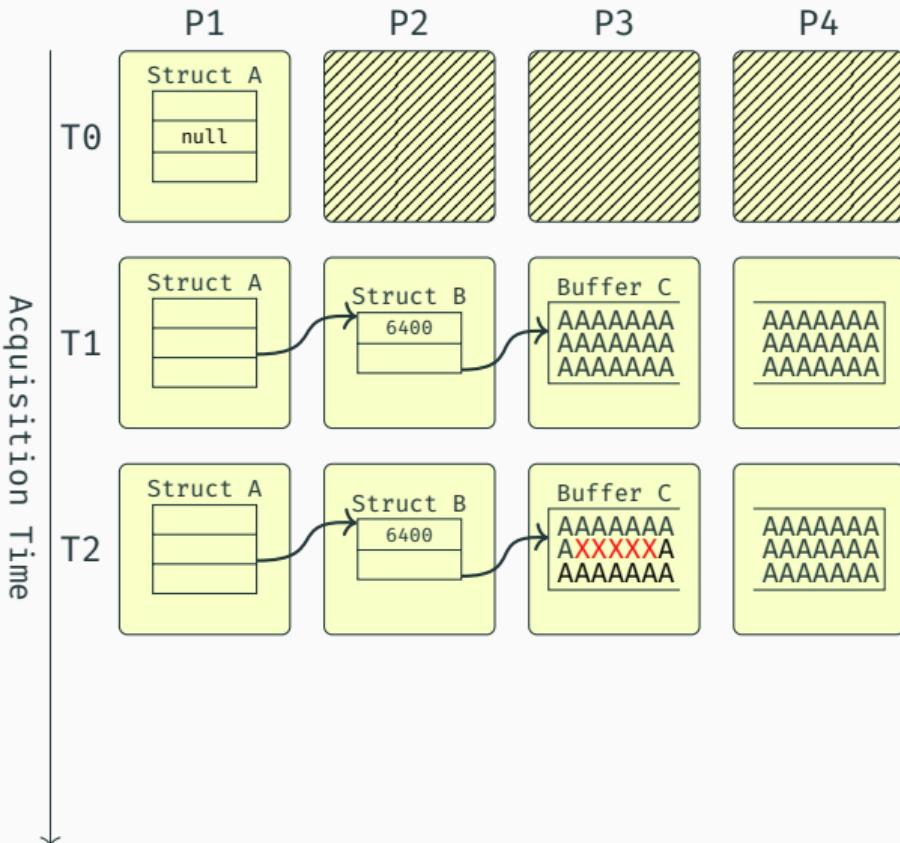
# Types of Inconsistency



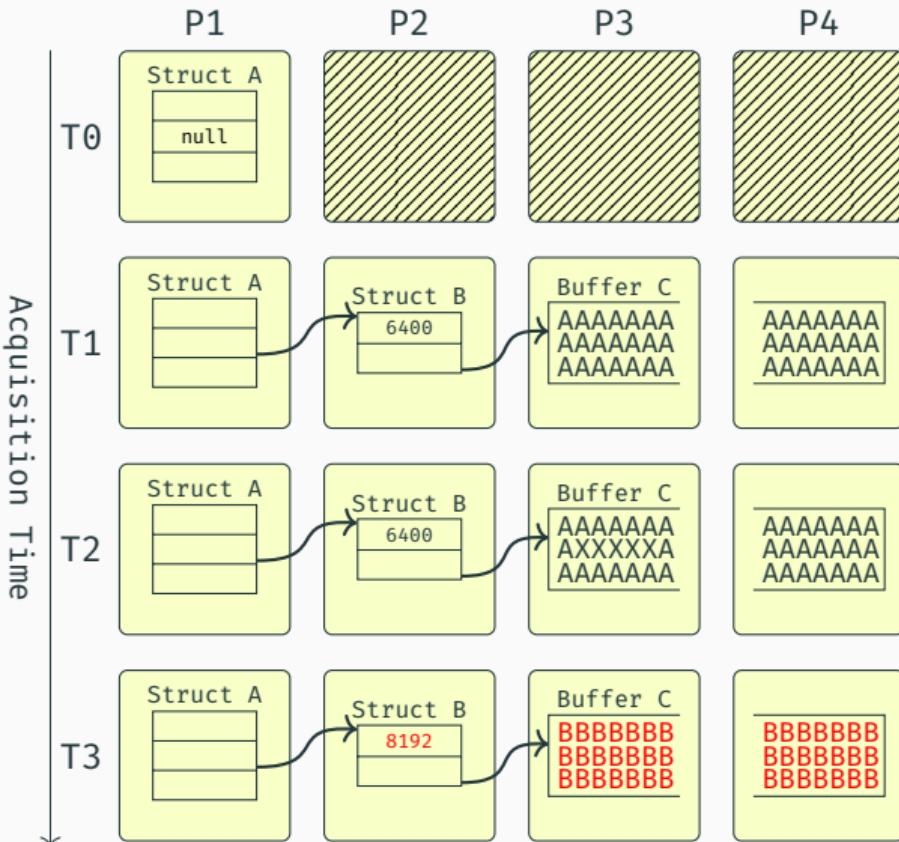
# Types of Inconsistency



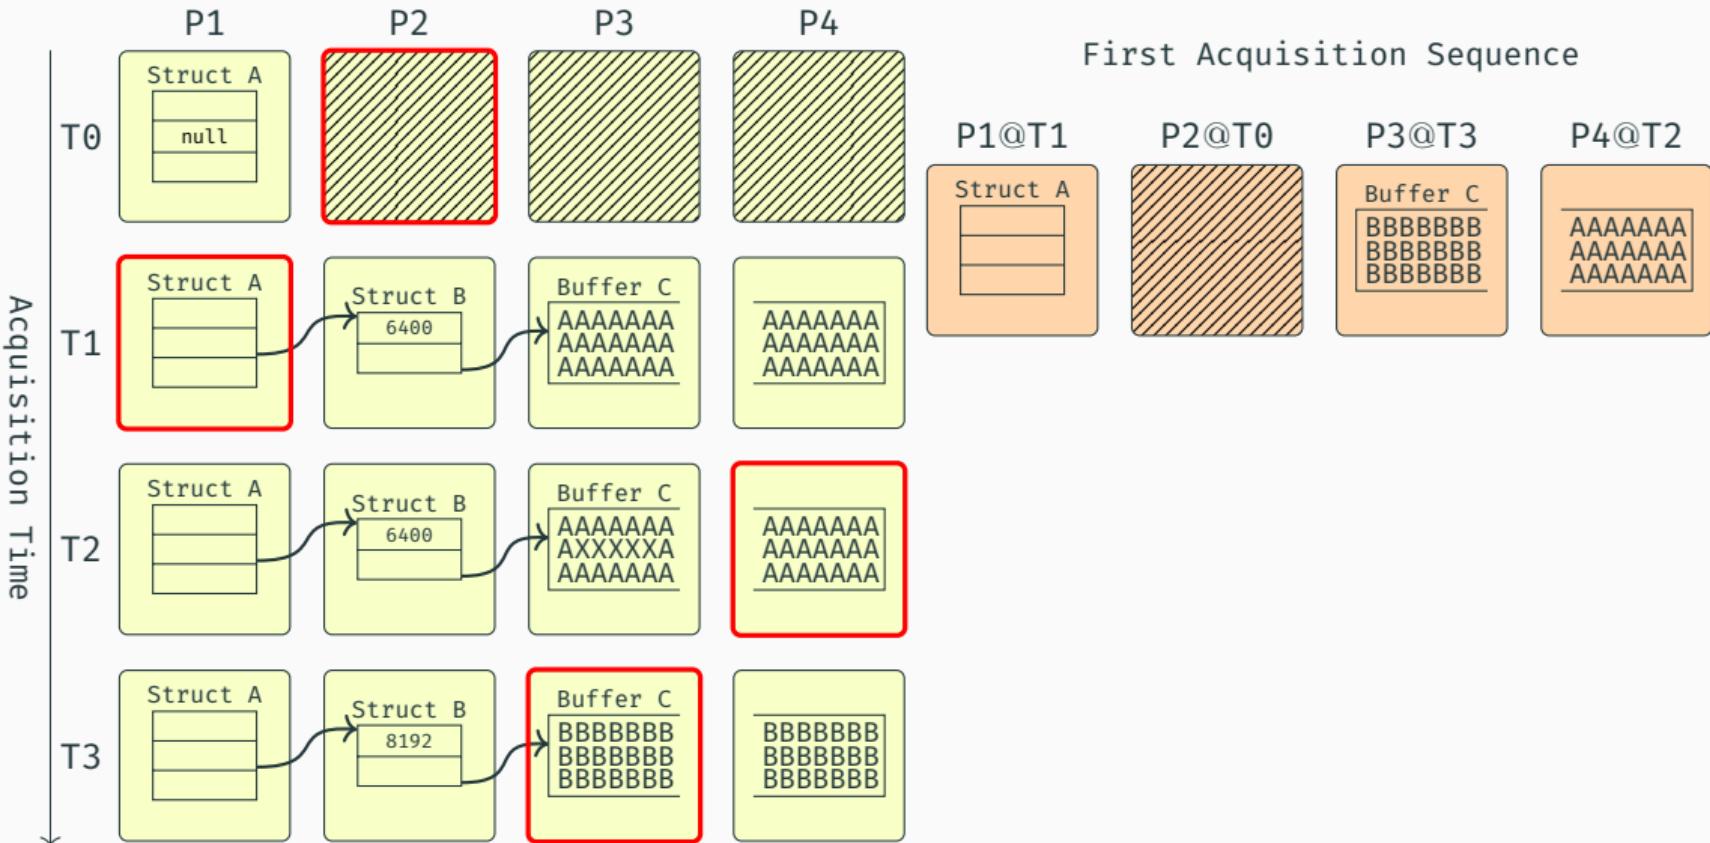
# Types of Inconsistency



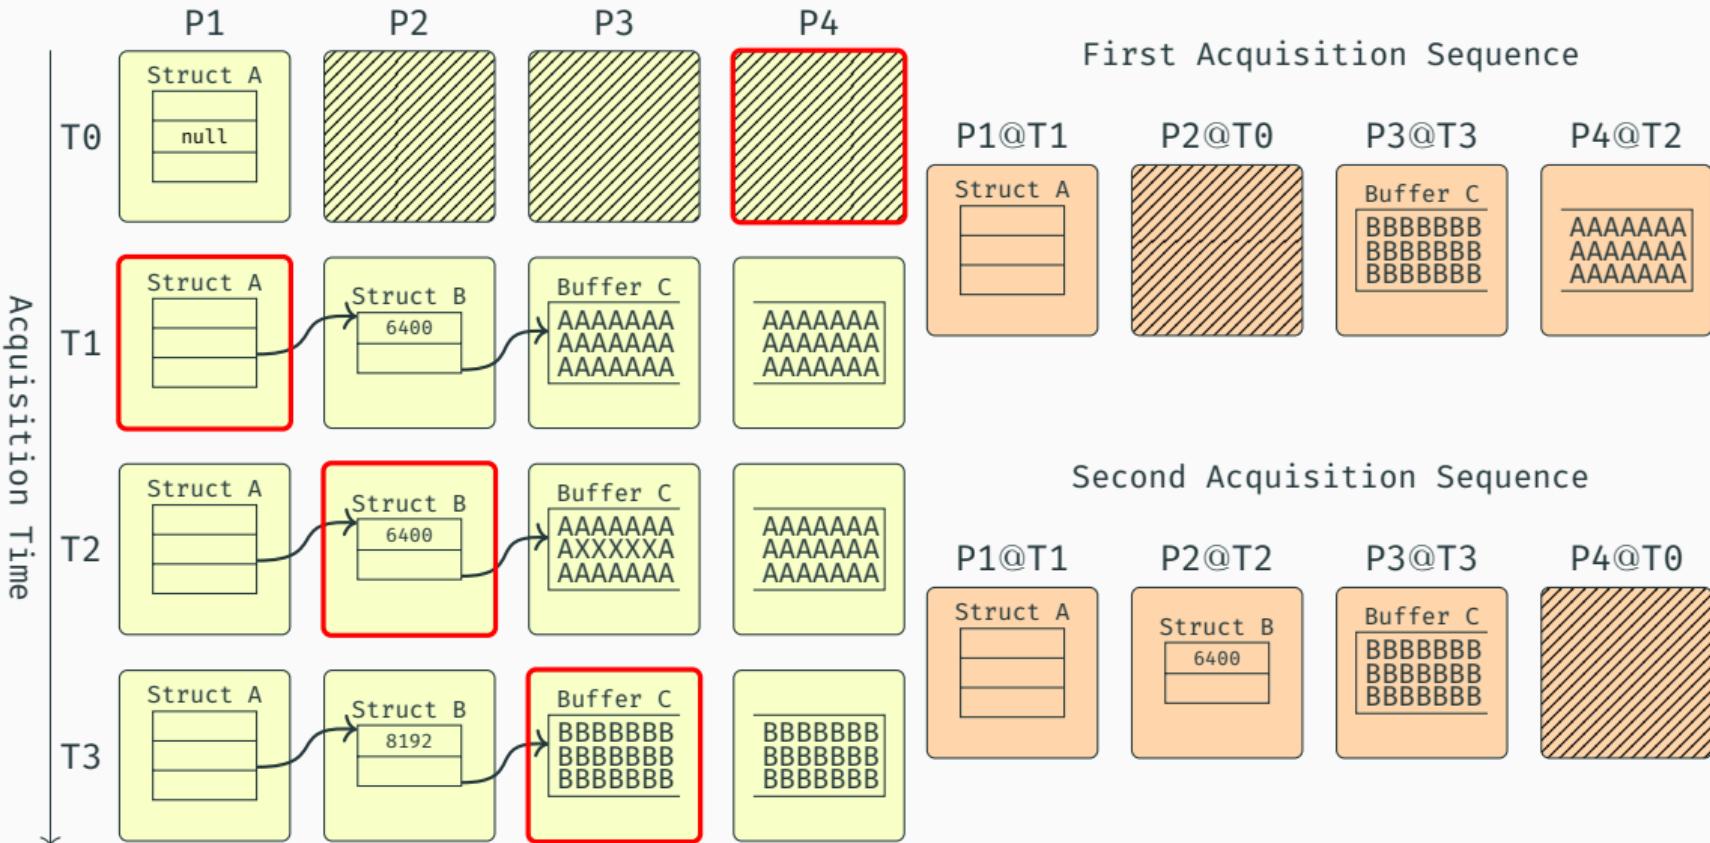
# Types of Inconsistency



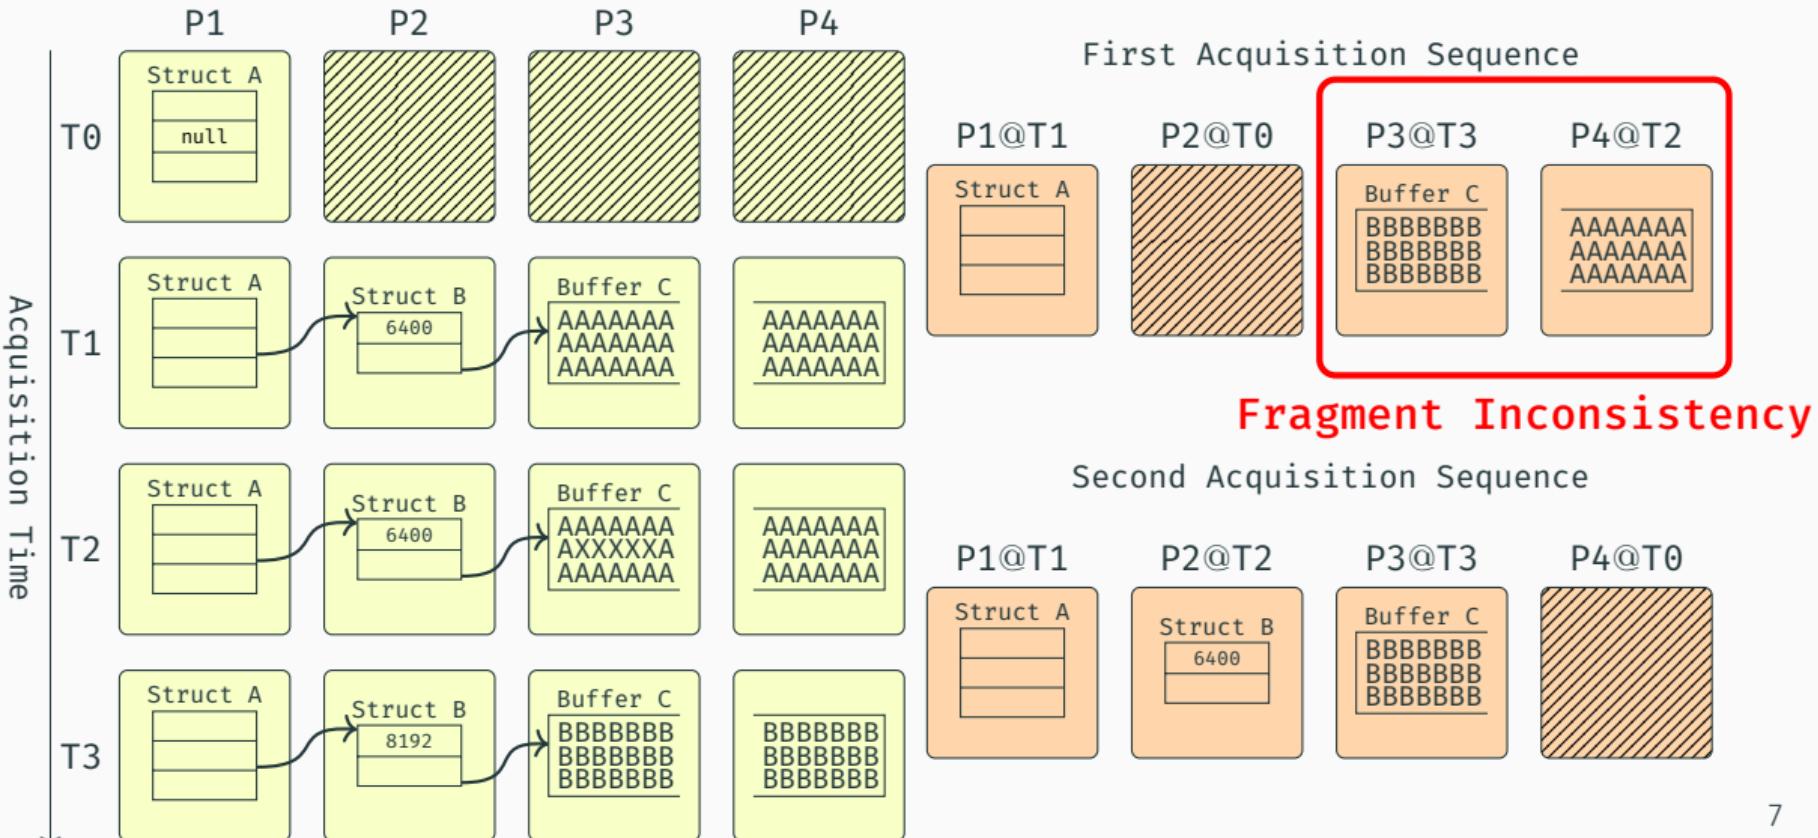
# Types of Inconsistency



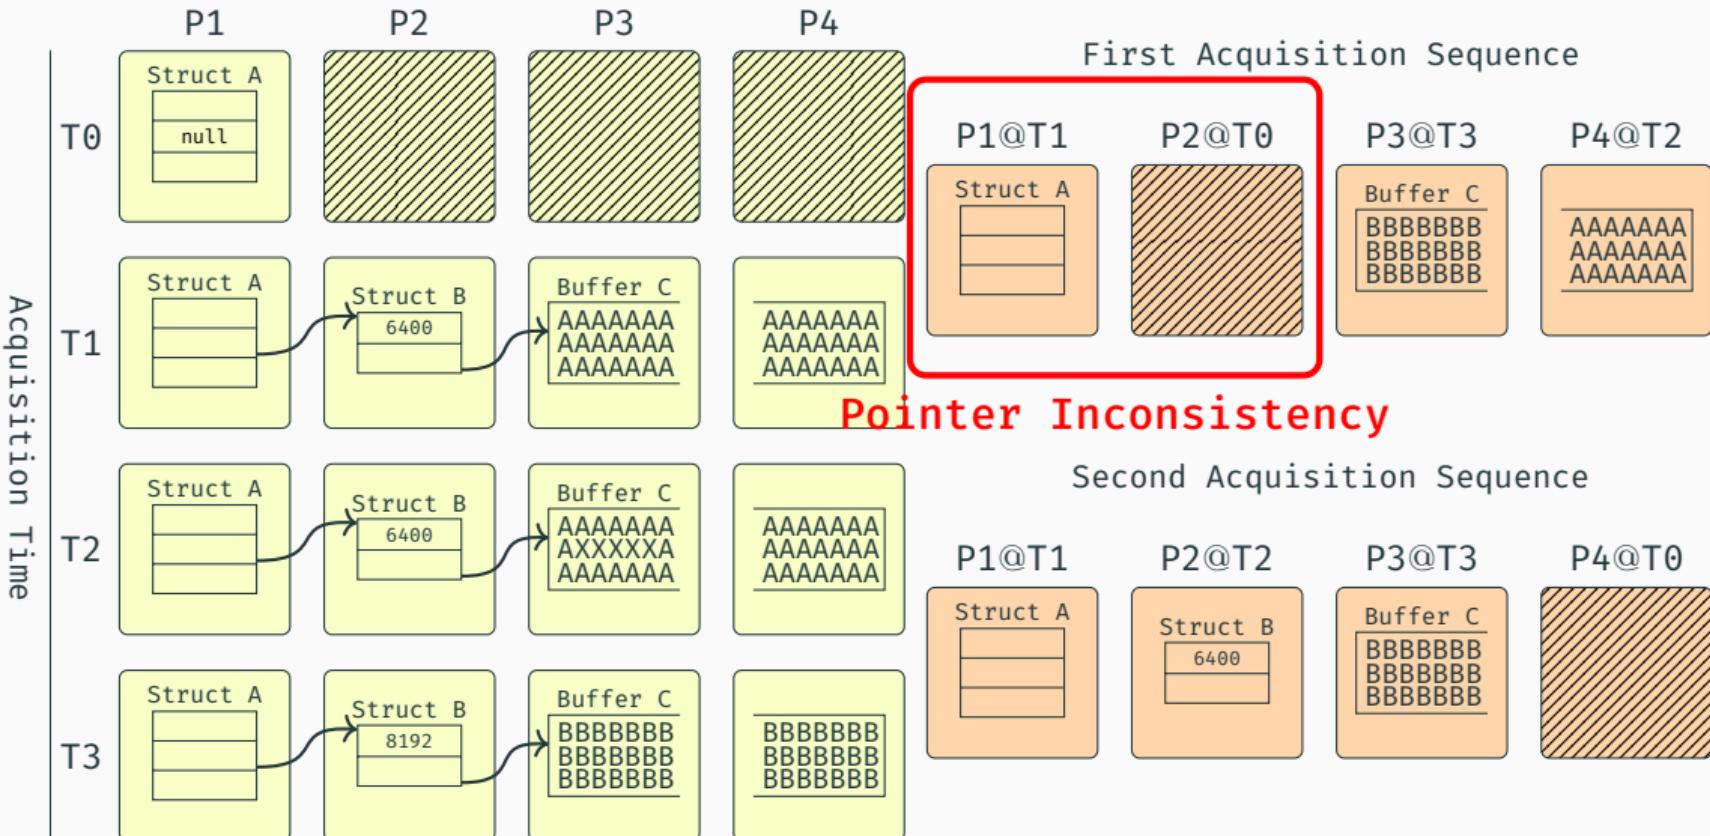
# Types of Inconsistency



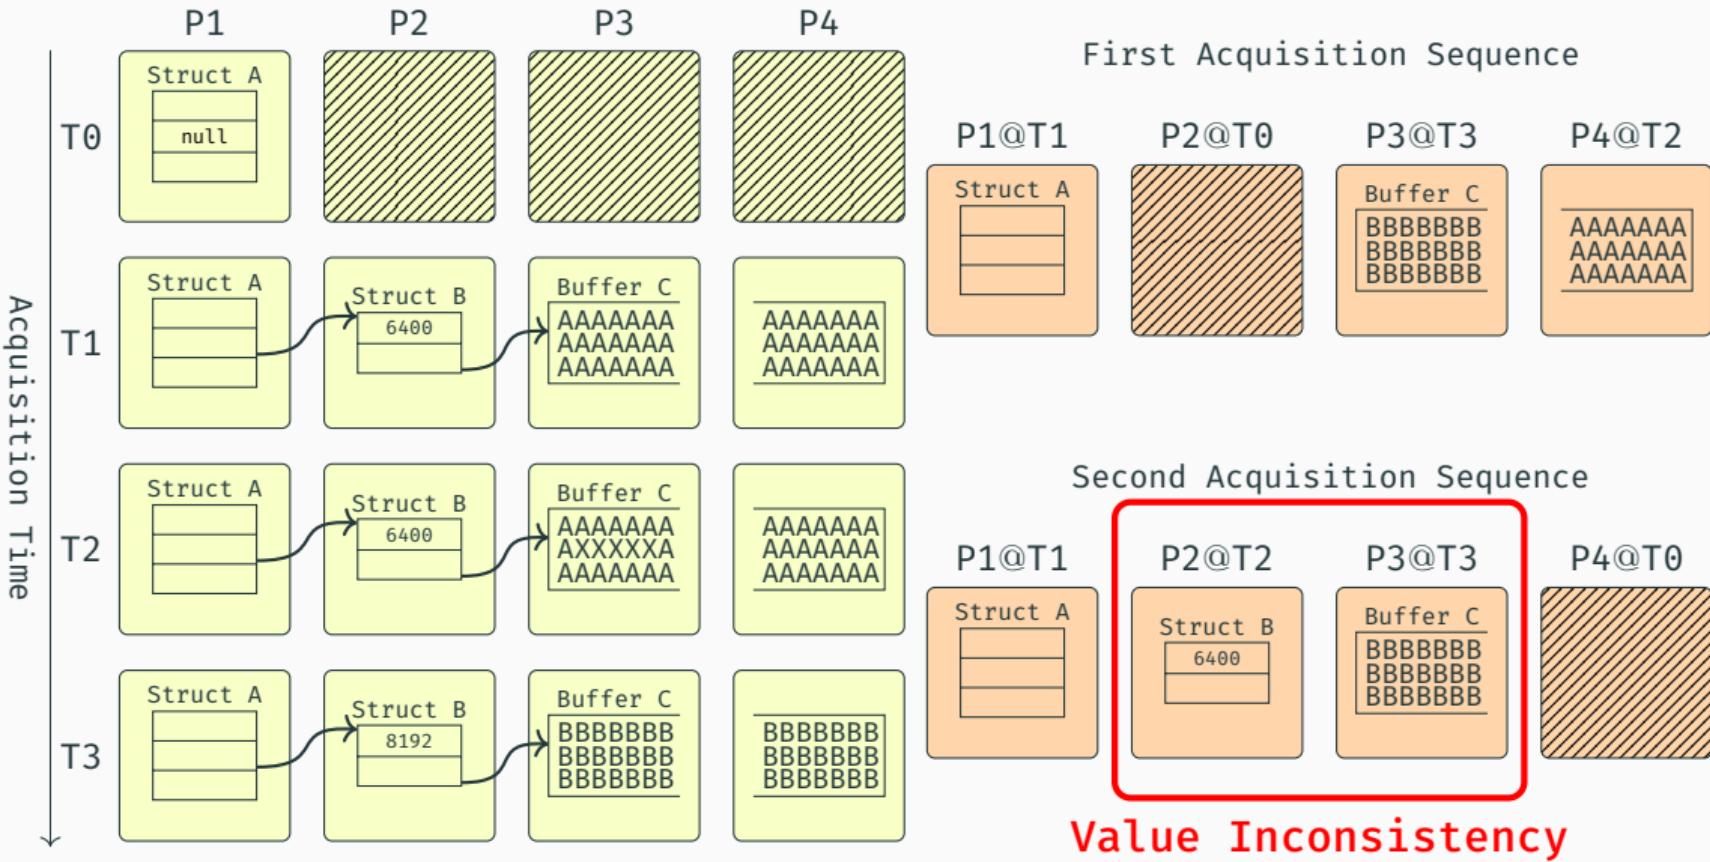
# Types of Inconsistency



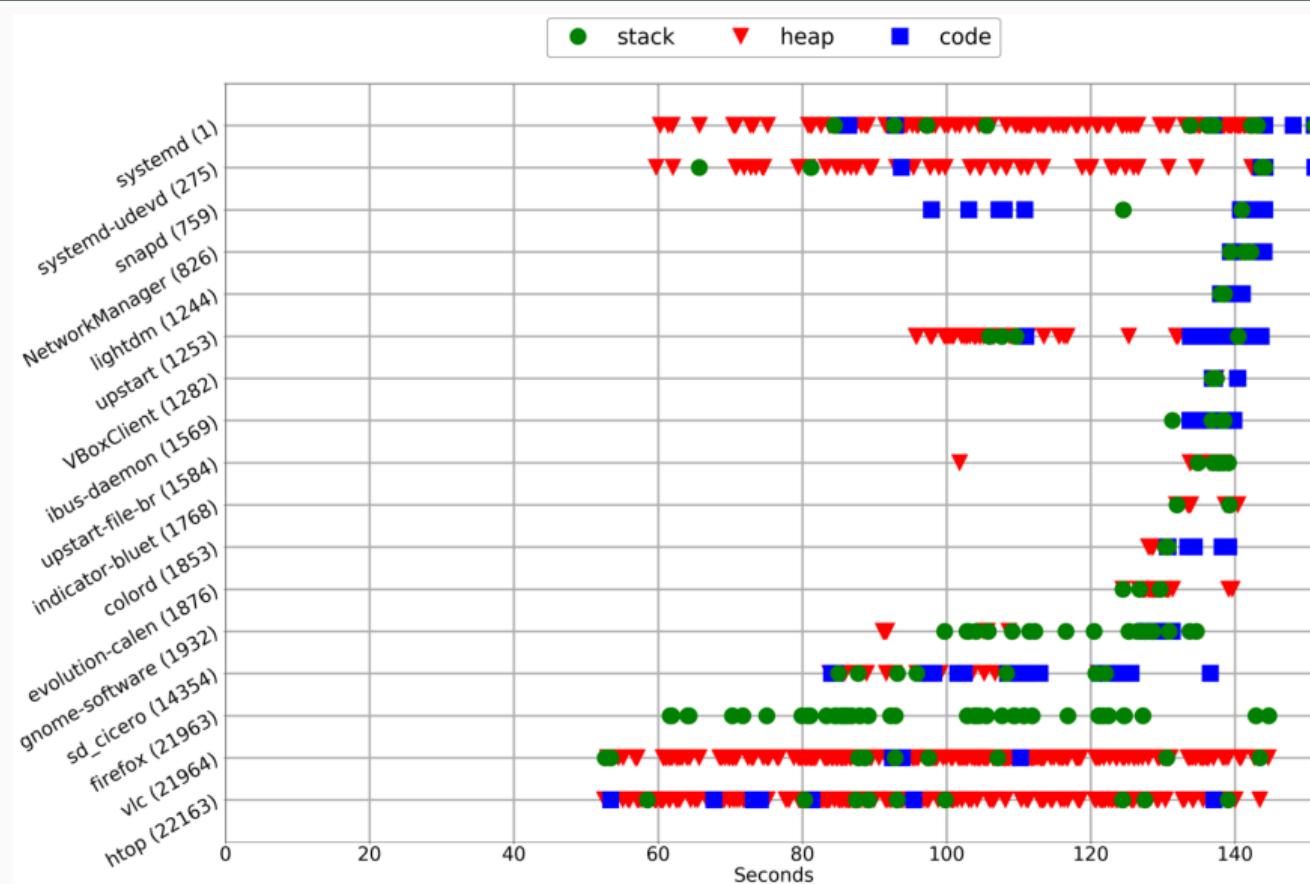
# Types of Inconsistency



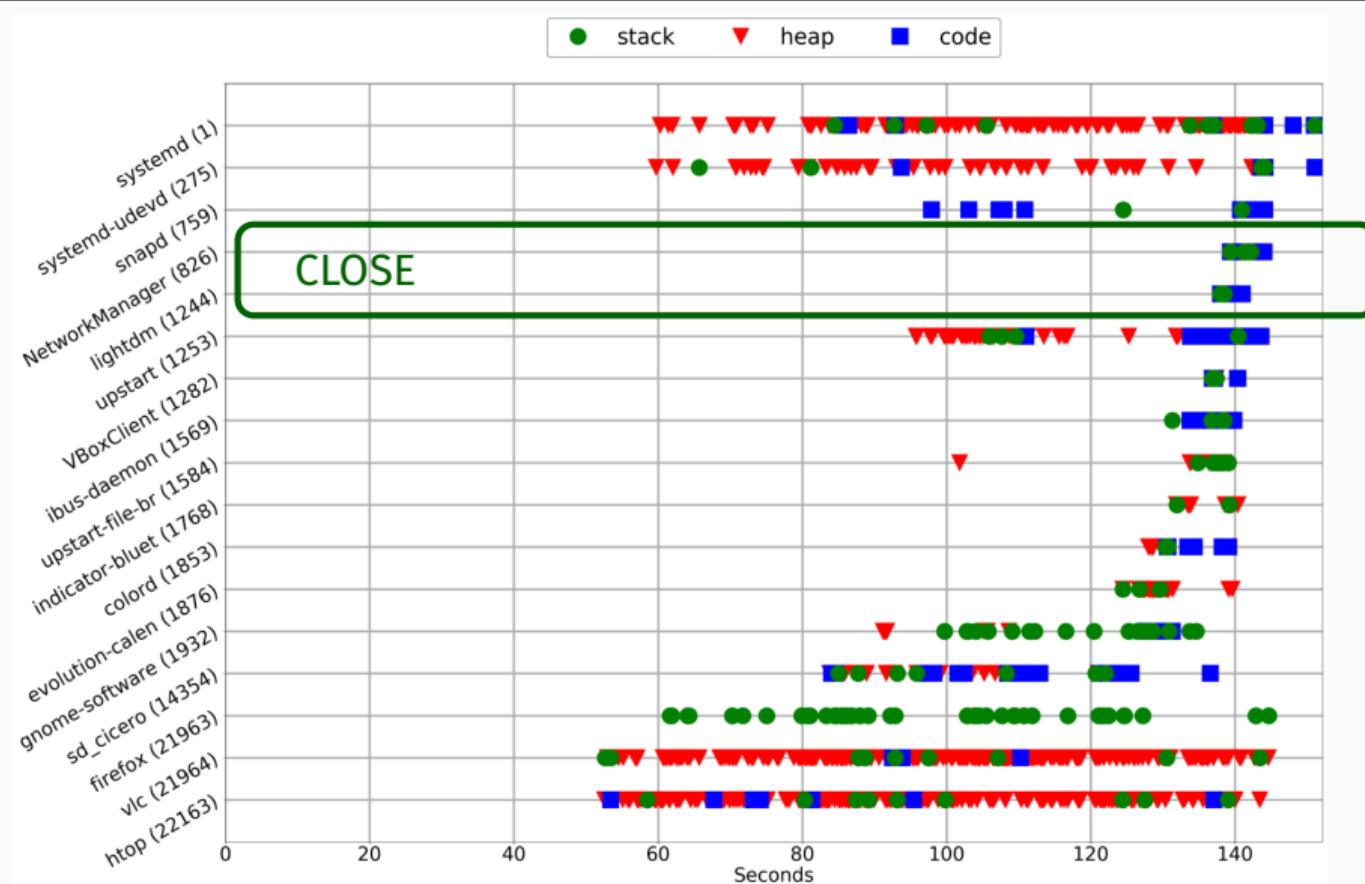
# Types of Inconsistency



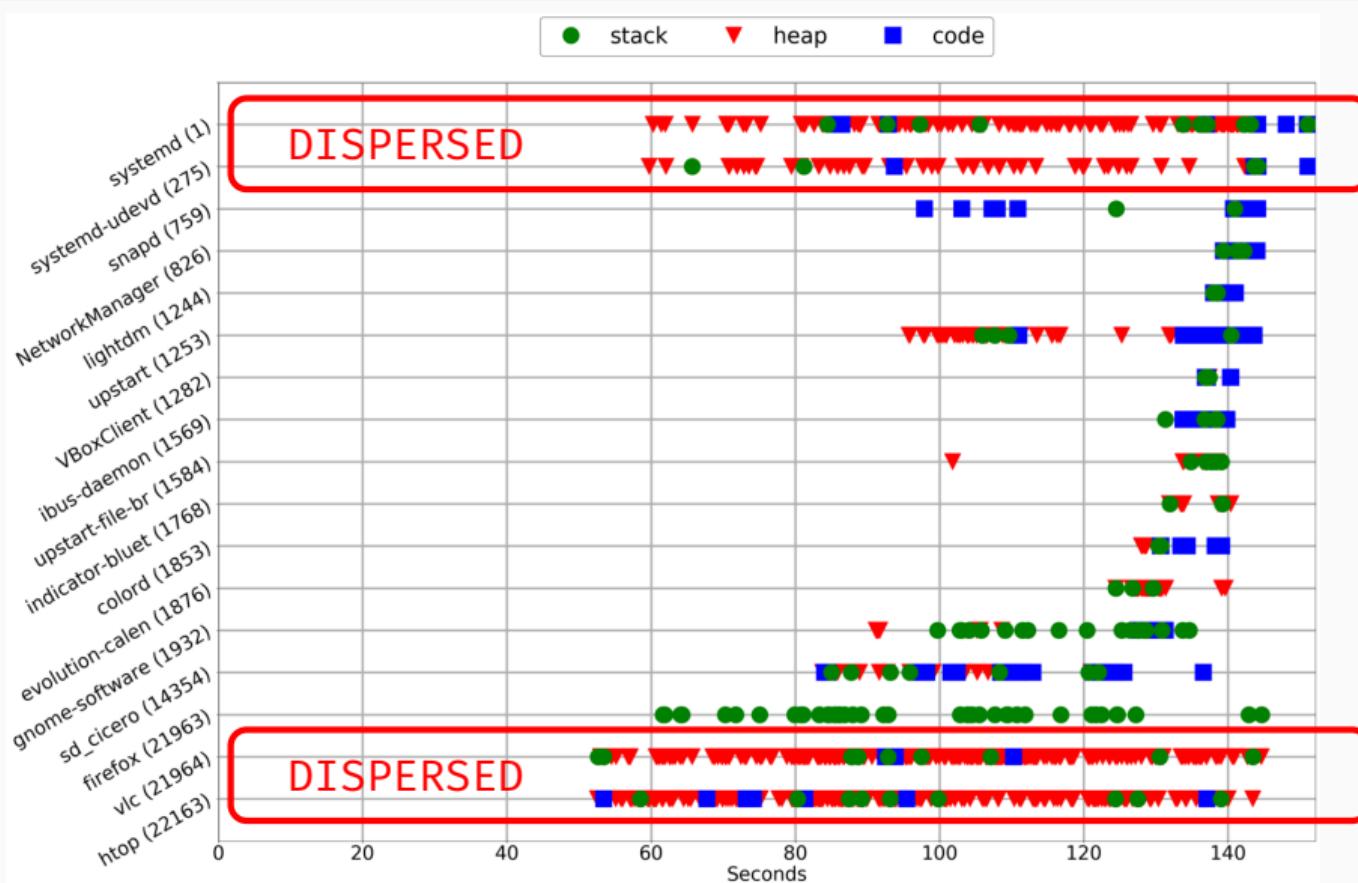
# Impact Estimation - Fragmentation



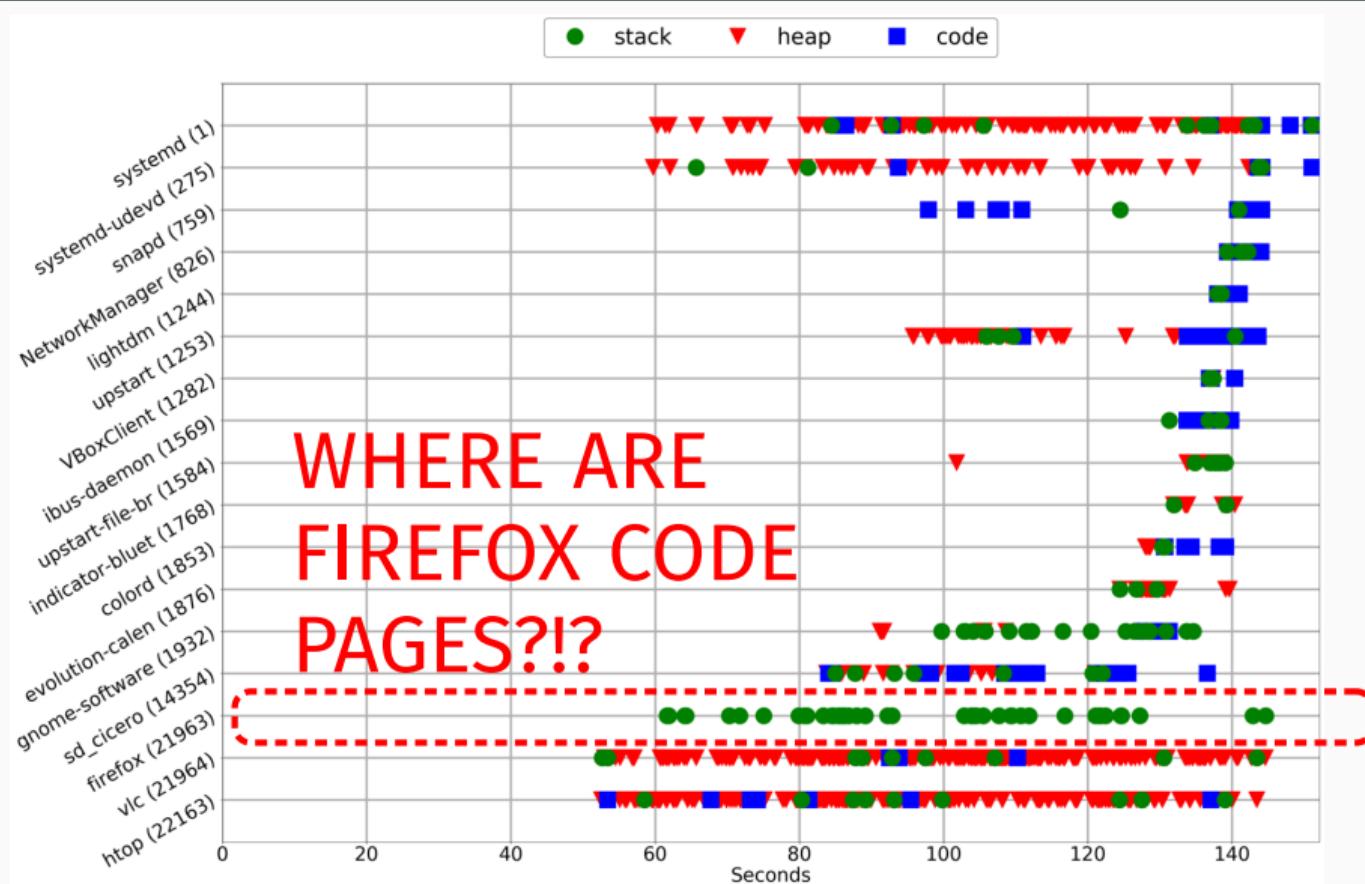
# Impact Estimation - Fragmentation



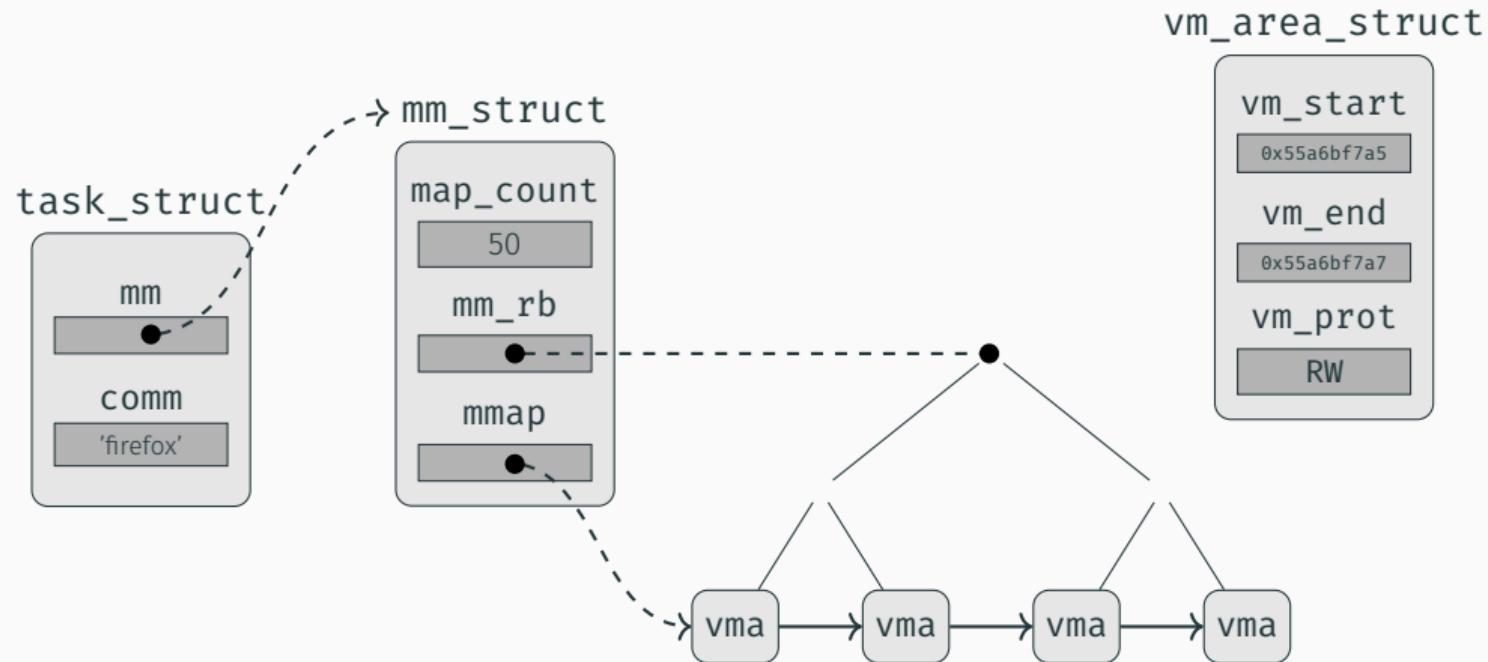
# Impact Estimation - Fragmentation



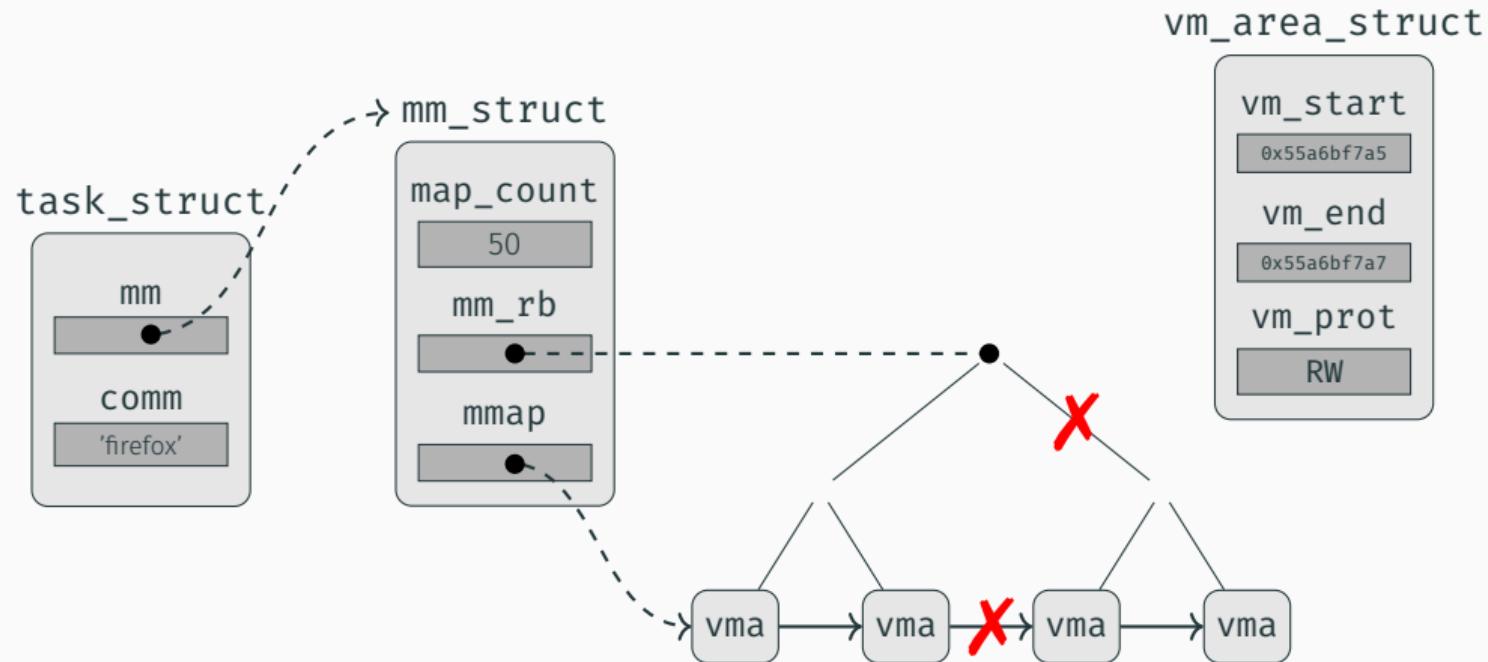
# Impact Estimation - Fragmentation



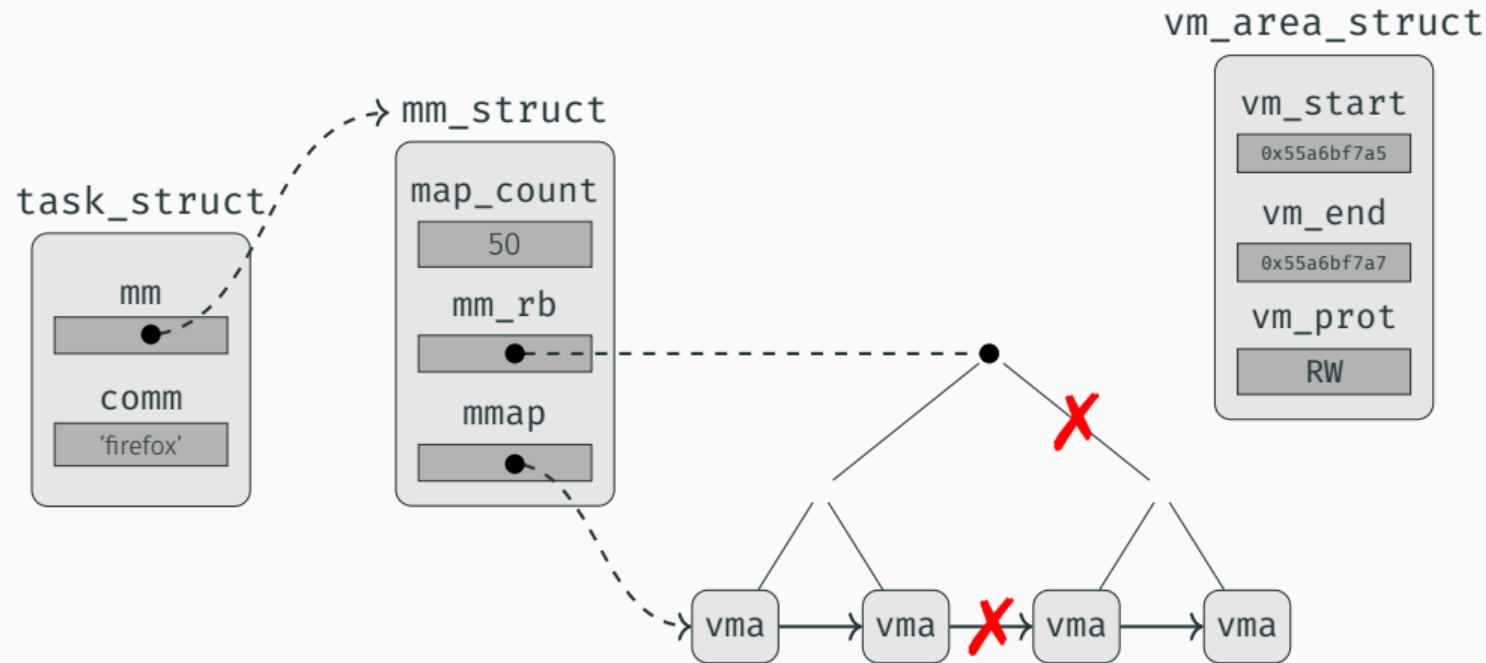
# Kernel-Space Integrity



# Kernel-Space Integrity



# Kernel-Space Integrity



Volatility Plugin: `map_count == list_len(mmap)`  
`map_count == tree_len(mm_rb)`

## Kernel-Space Integrity

|               | Scenario 1<br>(Firefox) | Scenario 2<br>(Apache) | Scenario 3<br>(Malware) |
|---------------|-------------------------|------------------------|-------------------------|
| List mismatch | 100%                    | 71%                    | 80%                     |
| Tree mismatch | 100%                    | 73%                    | 80%                     |
| Total         | 100%                    | 78%                    | 80%                     |

## Kernel-Space Integrity

|               | Scenario 1<br>(Firefox) | Scenario 2<br>(Apache) | Scenario 3<br>(Malware) |
|---------------|-------------------------|------------------------|-------------------------|
| List mismatch | 100%                    | 71%                    | 80%                     |
| Tree mismatch | 100%                    | 73%                    | 80%                     |
| Total         | 100%                    | 78%                    | 80%                     |

Is this actually a problem?

# Kernel-Space Integrity

|               | Scenario 1<br>(Firefox) | Scenario 2<br>(Apache) | Scenario 3<br>(Malware) |
|---------------|-------------------------|------------------------|-------------------------|
| List mismatch | 100%                    | 71%                    | 80%                     |
| Tree mismatch | 100%                    | 73%                    | 80%                     |
| Total         | 100%                    | 78%                    | 80%                     |

Is this actually a problem?

- List → Firefox stack and code **never** present
- Tree → Firefox stack present **10%**, code present **30%**

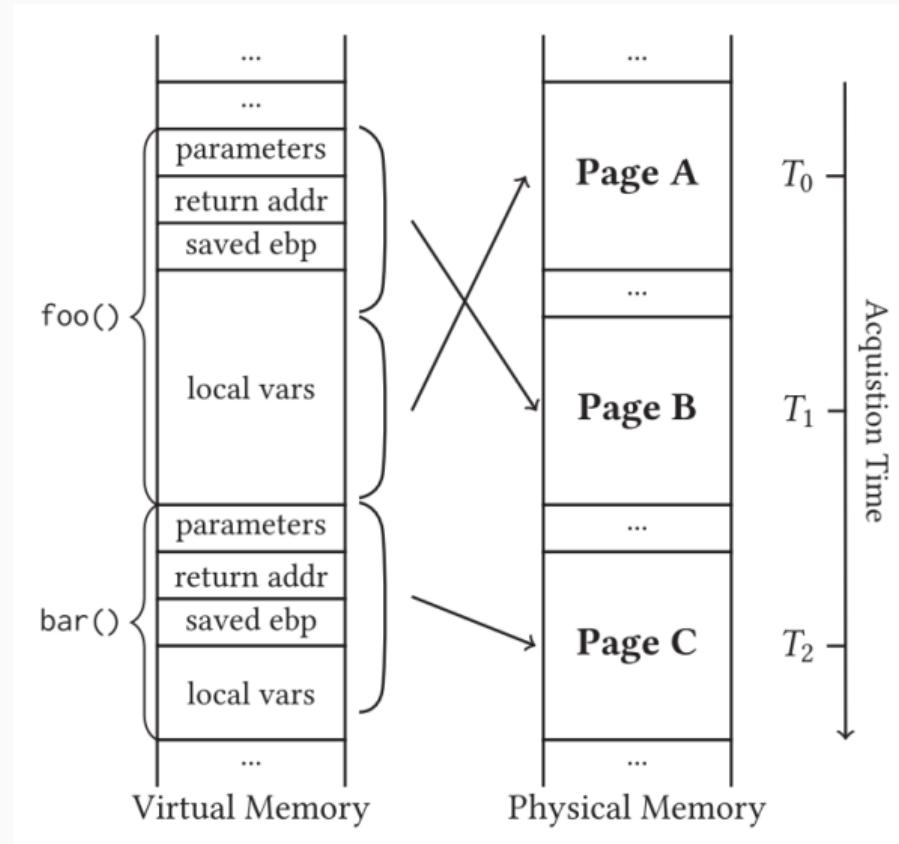
# Kernel-Space Integrity

|               | Scenario 1<br>(Firefox) | Scenario 2<br>(Apache) | Scenario 3<br>(Malware) |
|---------------|-------------------------|------------------------|-------------------------|
| List mismatch | 100%                    | 71%                    | 80%                     |
| Tree mismatch | 100%                    | 73%                    | 80%                     |
| Total         | 100%                    | 78%                    | 80%                     |

Is this actually a problem?

- List → Firefox stack and code **never** present
- Tree → Firefox stack present **10%**, code present **30%**
- Key recovery for WannaCry and NotPetya

# User-Space Integrity



# User-Space Integrity

|                            | $D_1$ | $D_2$ | $D_3$ | $D_4$ | $D_5$ | $D_6$ | $D_7$ | $D_8$ | $D_9$ | $D_{10}$ |
|----------------------------|-------|-------|-------|-------|-------|-------|-------|-------|-------|----------|
| Frames                     | -     | 6     | -     | 6     | 8     | -     | -     | -     | 6     | -        |
| Physical Pages             | 4     | 5     | 5     | 4     | 4     | 5     | 4     | 4     | 5     | 5        |
| Acquisition Time (s)       | 3.2   | 30.0  | 37.8  | 31.0  | 0.25  | 26.0  | 28.6  | 1.0   | 27.6  | 39.9     |
| rbp delta (s)              | 7.7   | 38.8  | 49.6  | 41.7  | 7.3   | 43.4  | 4.3   | 4.0   | 15.1  | 5.64     |
| Corrupted (registers)      | ✓     | -     | ✓     | -     | -     | -     | ✓     | ✓     | -     | ✓        |
| Corrupted (frame pointers) | -     | -     | -     | -     | -     | -     | ✓     | -     | -     | -        |
| Inconsistent data          | N/A   | ✓     | N/A   | ✓     | -     | N/A   | N/A   | N/A   | ✓     | N/A      |

## User-Space Integrity

|                            | $D_1$ | $D_2$ | $D_3$ | $D_4$ | $D_5$ | $D_6$ | $D_7$ | $D_8$ | $D_9$ | $D_{10}$ |
|----------------------------|-------|-------|-------|-------|-------|-------|-------|-------|-------|----------|
| Frames                     | -     | 6     | -     | 6     | 8     | -     | -     | -     | 6     | -        |
| Physical Pages             | 4     | 5     | 5     | 4     | 4     | 5     | 4     | 4     | 5     | 5        |
| Acquisition Time (s)       | 3.2   | 30.0  | 37.8  | 31.0  | 0.25  | 26.0  | 28.6  | 1.0   | 27.6  | 39.9     |
| rbp delta (s)              | 7.7   | 38.8  | 49.6  | 41.7  | 7.3   | 43.4  | 4.3   | 4.0   | 15.1  | 5.64     |
| Corrupted (registers)      | ✓     | -     | ✓     | -     | -     | -     | ✓     | ✓     | -     | ✓        |
| Corrupted (frame pointers) | -     | -     | -     | -     | -     | -     | ✓     | -     | -     | -        |
| Inconsistent data          | N/A   | ✓     | N/A   | ✓     | -     | N/A   | N/A   | N/A   | ✓     | N/A      |

Is this actually a problem?

# User-Space Integrity

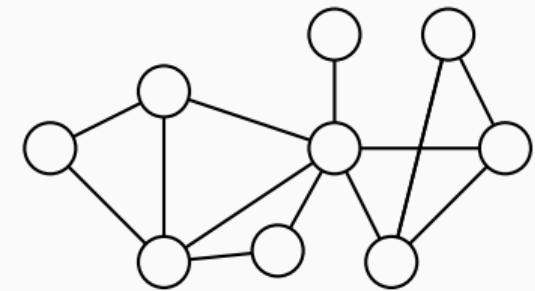
|                            | $D_1$ | $D_2$ | $D_3$ | $D_4$ | $D_5$ | $D_6$ | $D_7$ | $D_8$ | $D_9$ | $D_{10}$ |
|----------------------------|-------|-------|-------|-------|-------|-------|-------|-------|-------|----------|
| Frames                     | -     | 6     | -     | 6     | 8     | -     | -     | -     | 6     | -        |
| Physical Pages             | 4     | 5     | 5     | 4     | 4     | 5     | 4     | 4     | 5     | 5        |
| Acquisition Time (s)       | 3.2   | 30.0  | 37.8  | 31.0  | 0.25  | 26.0  | 28.6  | 1.0   | 27.6  | 39.9     |
| rbp delta (s)              | 7.7   | 38.8  | 49.6  | 41.7  | 7.3   | 43.4  | 4.3   | 4.0   | 15.1  | 5.64     |
| Corrupted (registers)      | ✓     | -     | ✓     | -     | -     | -     | ✓     | ✓     | -     | ✓        |
| Corrupted (frame pointers) | -     | -     | -     | -     | -     | -     | ✓     | -     | -     | -        |
| Inconsistent data          | N/A   | ✓     | N/A   | ✓     | -     | N/A   | N/A   | N/A   | ✓     | N/A      |

Is this actually a problem?

- Dissecting the user space process heap (DFRWS 2017)
- Building stack traces from memory dump of Windows x64 (DFRWS 2018)
- Chrome Ragamuffin (Volatility plugin for Chrome)

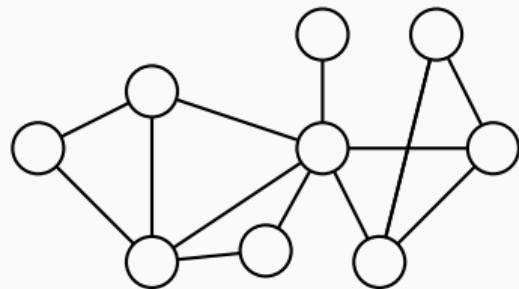
Can we study smear in a more generic way?

Build a graph of  
kernel structures

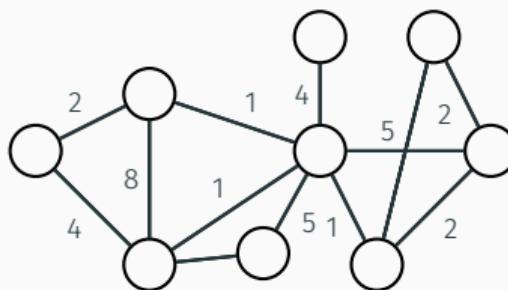


Can we study smear in a more generic way?

Build a graph of kernel structures

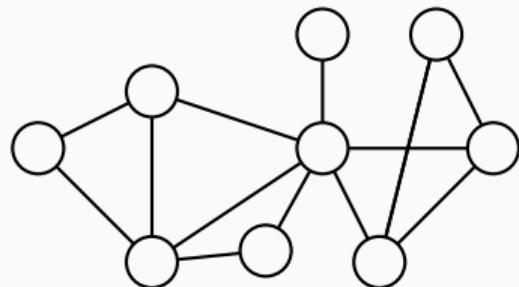


Define metrics to evaluate analyses

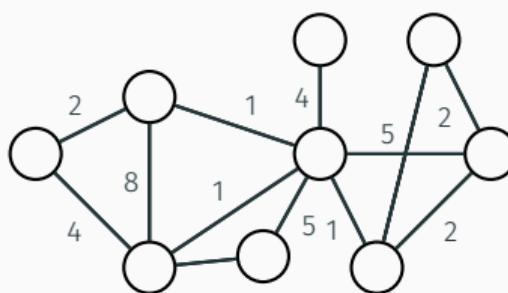


Can we study smear in a more generic way?

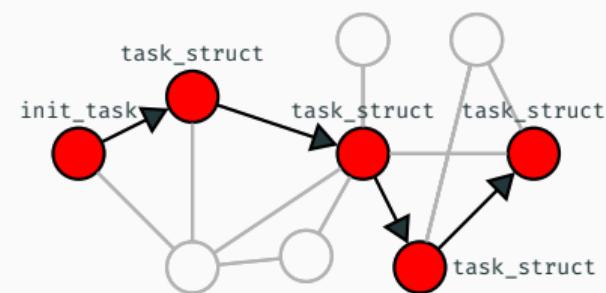
Build a graph of kernel structures



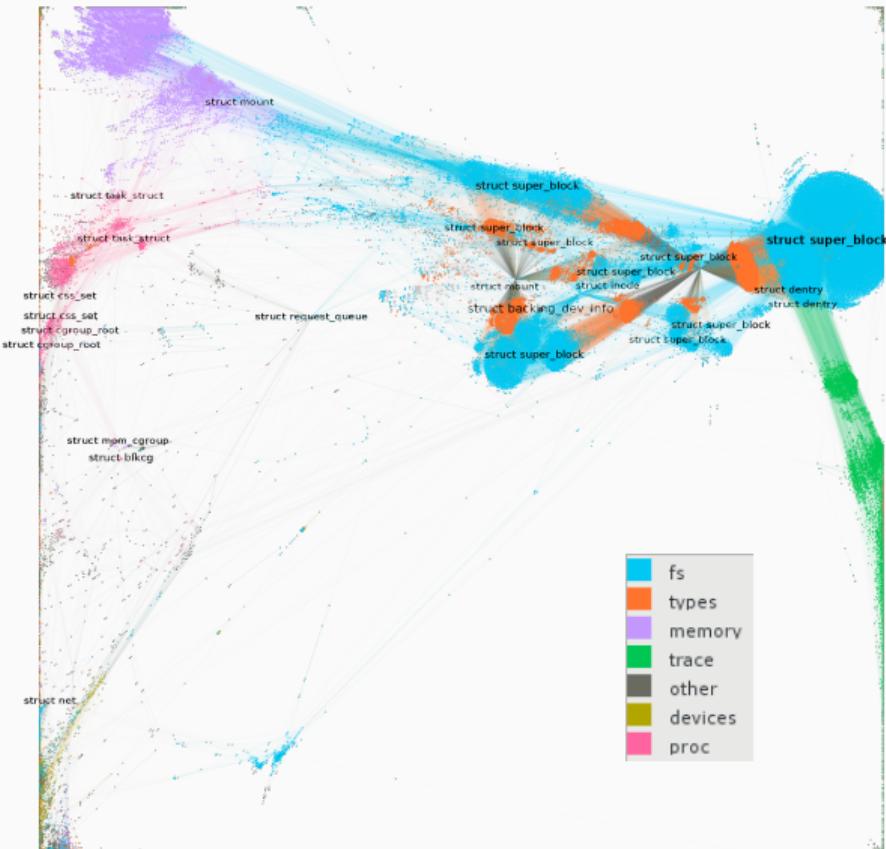
Define metrics to evaluate analyses



Study analyses as paths on the graph



# The Graph

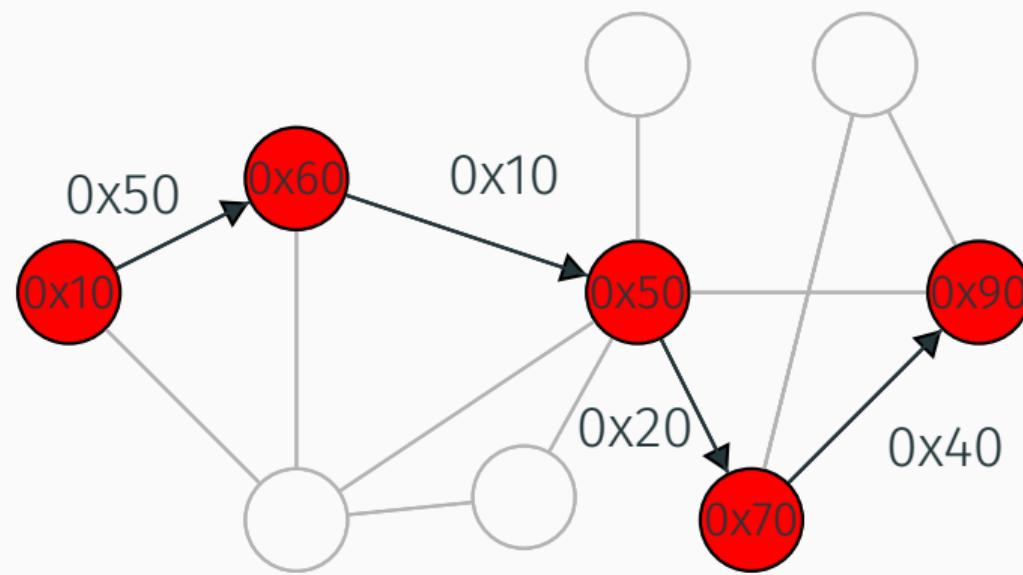


- 100k Structures (Nodes)
- 840k Pointers (Edges)

## Proposed Metrics

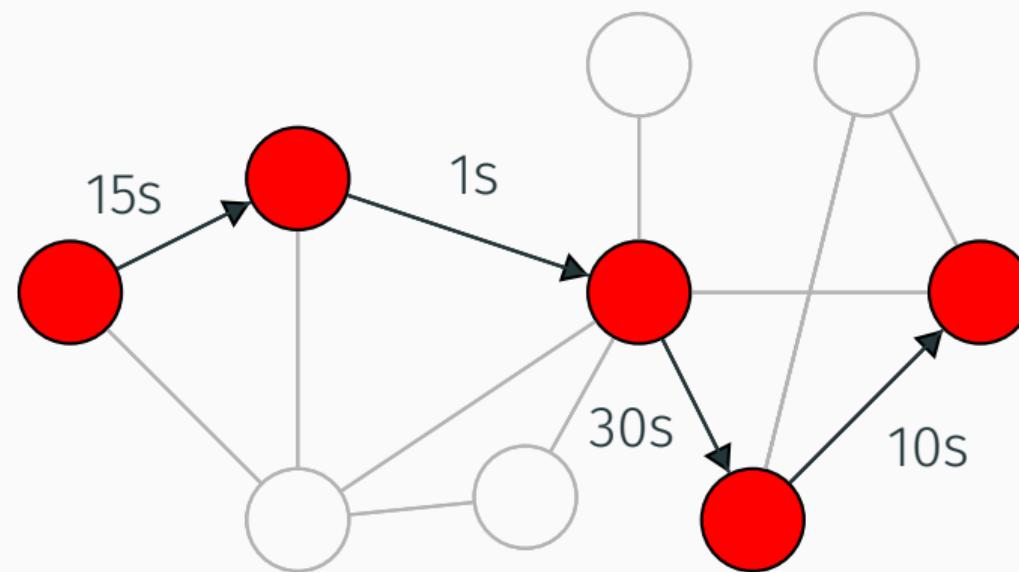
- Atomicity
- Stability
- Consistency

**Atomicity:** distance in memory between two connected structures

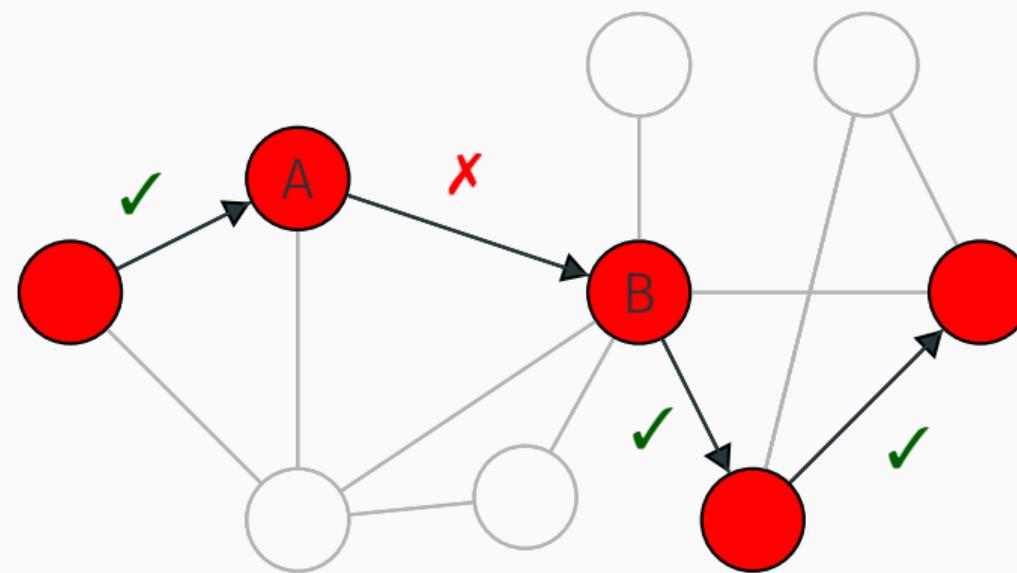


**Stability:** how long an edge remains stable in a running machine

- 25 snapshots at [0s, 1s, 5s, ..., 3h]



## Consistency: Atomicity + Stability



## Evaluation of Current Analyses

| Volatility Plugin   | # Nodes |  |  |
|---------------------|---------|--|--|
| linux_arp           | 13      |  |  |
| linux_check_creds   | 248     |  |  |
| linux_check_modules | 151     |  |  |
| linux_check_tty     | 13      |  |  |
| linux_find_file     | 14955   |  |  |
| linux_ifconfig      | 12      |  |  |
| linux_lsmod         | 12      |  |  |
| linux_lsof          | 821     |  |  |
| linux_mount         | 495     |  |  |
| linux_pidhashtable  | 469     |  |  |
| linux_proc_maps     | 4722    |  |  |
| linux_pslist        | 124     |  |  |

## Evaluation of Current Analyses

| Volatility Plugin   | # Nodes | Stability (s) |
|---------------------|---------|---------------|
| linux_arp           | 13      | 12,000        |
| linux_check_creds   | 248     | 2             |
| linux_check_modules | 151     | 700           |
| linux_check_tty     | 13      | 30            |
| linux_find_file     | 14955   | 0             |
| linux_ifconfig      | 12      | 12,000        |
| linux_lsmod         | 12      | 700           |
| linux_lsof          | 821     | 0             |
| linux_mount         | 495     | 10            |
| linux_pidhashtable  | 469     | 30            |
| linux_proc_maps     | 4722    | 0             |
| linux_pslist        | 124     | 30            |

Stability: 3 paths **never** changed in over 3 hours  
11 paths **changed** in less than 1 minute

## Evaluation of Current Analyses

| Volatility Plugin   | # Nodes | Stability (s) | Consistency |      |
|---------------------|---------|---------------|-------------|------|
|                     |         |               | Fast        | Slow |
| linux_arp           | 13      | 12,000        | ✓           | ✓    |
| linux_check_creds   | 248     | 2             | ✓           | ✓    |
| linux_check_modules | 151     | 700           | ✓           | ✓    |
| linux_check_tty     | 13      | 30            | ✓           | ✓    |
| linux_find_file     | 14955   | 0             | ✗           | ✗    |
| linux_ifconfig      | 12      | 12,000        | ✓           | ✓    |
| linux_lsmod         | 12      | 700           | ✓           | ✓    |
| linux_lsof          | 821     | 0             | ✗           | ✗    |
| linux_mount         | 495     | 10            | ✓           | ✗    |
| linux_pidhashtable  | 469     | 30            | ✓           | ✗    |
| linux_proc_maps     | 4722    | 0             | ✗           | ✗    |
| linux_pslist        | 124     | 30            | ✓           | ✓    |

Consistency: 5 inconsistent plugins when fast acquisition  
7 inconsistent plugins when slow acquisition

## Solutions

---

## A New Temporal Dimension - Recording Time

- Given a physical page we must be able to tell *when* it was acquired!
- Modified LiME to record timing information
- Overhead:
  - Every  $100\mu s \rightarrow 0.7\%$
  - Every page  $\rightarrow 2.4\%$

## A New Temporal Dimension - Time Analysis

- Transparently add the timing information to Volatility
- Intercept object creation to create a *timeline*:

---

```
./vol.py -f dump.raw --profile=... --pagetime pslist  
<original pslist output>
```

Accessed physical pages: 171  
Acquisition time window: 72s

[XX-----XxX---xXXX--xX-xX---Xxx-xx-X-XxxX-XXX]

---

## Locality-Based Acquisition

- Every memory acquisition tool treats pages equally:
  - Independently if it is used by the OS
  - Independently if it contains forensics data
  - From lowest → highest physical address
- Can we do better?
- Why not acquiring forensics/interconnected data first, and then rest of memory?

## Locality-Based Acquisition

Two phases:

1. *Smart* dump:

- Process and module list
- For each process: page tables, memory mappings, open files, stack, heap, kernel stack..

2. Traditional acquisition of the remaining pages

## Locality-Based Acquisition

Two phases:

1. *Smart* dump:

- Process and module list
- For each process: page tables, memory mappings, open files, stack, heap, kernel stack..

2. Traditional acquisition of the remaining pages

### Impact

- Negligible overhead in time and memory footprint

## Locality-Based Acquisition

Two phases:

1. *Smart* dump:

- Process and module list
- For each process: page tables, memory mappings, open files, stack, heap, kernel stack..

2. Traditional acquisition of the remaining pages

### Impact

- Negligible overhead in time and memory footprint
- No inconsistency in kernel and user space integrity tests!

# DEMO

More details on our papers:

- Introducing the Temporal Dimension to Memory Forensics (ACM TOPS 2019)
- Back to the Whiteboard: a Principled Approach for the Assessment and Design of Memory Forensic Techniques (USENIX 2019)

All the code and artifacts developed are open-source!

- [https://github.com/pagabuc/atomicity\\_tops](https://github.com/pagabuc/atomicity_tops)
- <https://github.com/pagabuc/kernographer>

# Questions?

Twitter: @pagabuc

Email: pagani@eurecom.fr